

ISIT312 Big Data Management

Physical Data Warehouse Design

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Physical Data Warehouse Design

Outline

[Techniques for Physical Data Warehouse Design](#)

[Materialized View](#)

[Indexes for Data Warehouses](#)

[Evaluation of Star Queries](#)

[Data Warehouse Partitioning](#)

Techniques for Physical Data Warehouse Design

Materialized Views

- A view physically stored in the DB
- Typical problems: view update, view selection

Indexing

- Used in Data Warehouse together with materialized views
- Specific for Data Warehouse: bitmap and join indexes

Partitioning

- Divides the contents of a relation in several files
- Horizontal and vertical partitioning

Physical Data Warehouse Design

Outline

[Techniques for Physical Data Warehouse Design](#)

[Materialized View](#)

[Indexes for Data Warehouses](#)

[Evaluation of Star Queries](#)

[Data Warehouse Partitioning](#)

Materialized Views

Materialized view is a relational table that contains the rows that would be returned by the view definition - usually SELECT statement of SQL

If we consider relational views as stored queries then **materialized views** can be considered as stored results

Materialized views are created and used to reduce an amount of time needed to compute SELECT statements, for example join materialized views eliminate the needs to join the relational table

There are two ways how materialized view can be used:

- brute force method
- transparent query rewrite

In brute force method SQL is written to explicitly access the view

Transparent query rewrite method is applied when a query optimizer detects that a query can be computed against a materialized view instead of the source relational tables

Materialized Views

~~View maintenance~~ means that when base relational tables are updated, view must be updated too

Incremental view maintenance means that updated view is computed from the individual modifications to the relational tables and not from the entire relational tables

Creating **materialized view**

```
CREATE MATERIALIZED VIEW MV_ORDERS  
REFRESH ON COMMIT  
ENABLE QUERY REWRITE  
AS( SELECT O_ORDERKEY, O_CUSTKEY, O_TOTALPRICE, O_ORDERDATE  
FROM ORDERS  
WHERE O_ORDERDATE > TO_DATE('31-DEC-1986', 'DD-MON-YYYY') );
```

Creating materialized view

Direct access to **materialized view**

```
SELECT *  
FROM MV_ORDERS  
WHERE O_ORDERDATE = TO_DATE('01-JAN-1992', 'DD-MON-YYYY')
```

Direct access to materialized view

Materialized Views

Access to materialized view through query rewriting

Indirect access to materialized view through query rewriting

```
SELECT O_ORDERKEY, O_CUSTKEY, O_TOTALPRICE, O_ORDERDATE
FROM ORDERS
WHERE O_ORDERDATE > TO_DATE('31-DEC-1986', 'DD-MON-YYYY');
```

- The results from **EXPLAIN PLAN** statement

Query processing plan

PLAN_TABLE_OUTPUT

0	SELECT STATEMENT		108K	2539K	507	(1)	00:00:01
* 1	MAT_VIEW REWRITE ACCESS FULL	MV_ORDERS	108K	2539K	507	(1)	00:00:01

Predicate Information (identified by operation id):

```
1 - filter("MV_ORDERS"."O_ORDERDATE">TO_DATE(' 1986-12-31 00:00:00',
'syyyy-mm-dd hh24:mi:ss'))
```

Physical Data Warehouse Design

Outline

[Techniques for Physical Data Warehouse Design](#)

[Materialized View](#)

[Indexes for Data Warehouses](#)

[Evaluation of Star Queries](#)

[Data Warehouse Partitioning](#)

Indexes for Data Warehouses

An index provides a quick way to locate data of interest

Sample query

```
SELECT *  
FROM EMPLOYEE  
WHERE EmployeeKey = 007;
```

SELECT statement with equality condition in WHERE clause

With the help of an index over a column **EmployeeKey** (primary key in **EMPLOYEE** table), a single disk block access will suffice to answer the query

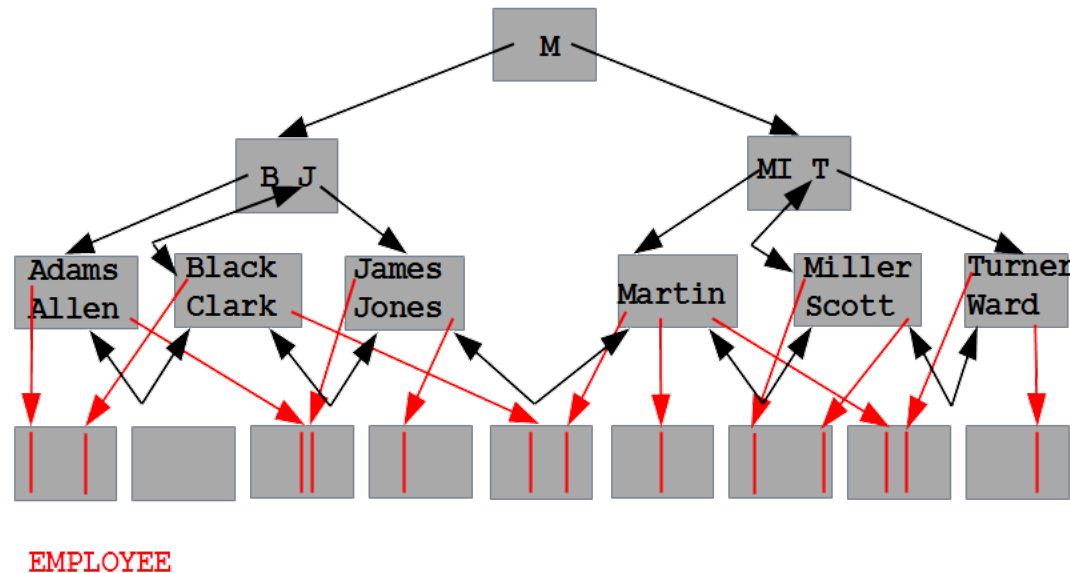
Without this index, we should perform a complete scan of table **EMPLOYEE**

Drawback: Almost every update on an indexed attribute also requires an index update

Too many indexes may degrade performance

Most popular indexing techniques in relational databases include **B*-trees** and **bitmap indexes**

B*-tree index implementation



B*-tree can be traversed either:

- vertically from root to leaf level of a tree
- horizontally either from left corner of leaf level to right corner of leaf level or the opposite
- vertically and later on horizontally either towards left lower corner or right lower corner of leaf level

Bitmap Indexes

ProductKey	ProductName	QuantityPerUnit	UnitPrice	Discontinued	CategoryKey
p1	prod1	25	60	No	c1
p2	prod2	45	60	Yes	c1
p3	prod3	50	75	No	c2
p4	prod4	50	100	Yes	c2
p5	prod5	50	120	No	c3
p6	prod6	70	110	Yes	c4

Product dimension table

	25	45	50	70
p1	1	0	0	0
p2	0	1	0	0
p3	0	0	1	0
p4	0	0	1	0
p5	0	0	1	0
p6	0	0	0	1

Bitmap index for attribute QuantityPerUnit

	60	75	100	110	120
p1	1	0	0	0	0
p2	1	0	0	0	0
p3	0	1	0	0	0
p4	0	0	1	0	0
p5	0	0	0	0	1
p6	0	0	0	1	0

Bitmap index for attribute UnitPrice

Bitmap Indexes: Example

Products having **between 45 and 55** pieces per unit, **and** with a **unit price between 100 and 200**

	45	50	OR1
p1	0	0	0
p2	1	0	1
p3	0	1	1
p4	0	1	1
p5	0	1	1
p6	0	0	0

OR for QuantityPerUnit

	100	110	120	OR2
p1	0	0	0	0
p2	0	0	0	0
p3	0	0	0	0
p4	1	0	0	1
p5	0	0	1	1
p6	0	1	0	1

OR for UnitPrice

	OR1	OR2	AND
p1	0	0	0
p2	1	0	0
p3	1	0	0
p4	1	1	1
p5	1	1	1
p6	0	1	0

AND operation

Indexes for Data Warehouses: Requirements

Symmetric partial match queries

- All dimensions of the cube should be symmetrically indexed, to be searched simultaneously

Indexing at multiple levels of aggregation

- Summary tables must be indexed in the same way as base nonaggregated tables

Efficient batch update

- The refreshing time of a data warehouse must be considered when designing the indexing schema

Sparse data

- Typically, only 20% of the cells in a data cube are nonempty
- The indexing schema must deal efficiently with sparse and nonsparse data

Physical Data Warehouse Design

Outline

[Techniques for Physical Data Warehouse Design](#)

[Materialized View](#)

[Indexes for Data Warehouses](#)

[Evaluation of Star Queries](#)

[Data Warehouse Partitioning](#)

Star Queries

Queries over star schemas are called **star queries**

Join the fact table with the dimension tables

A typical star query: total sales of discontinued products, by customer name and product name

```
SELECT ProductName, CustomerName, SUM(SalesAmount)
FROM Sales S, Customer C, Product P
WHERE S.CustomerKey = C.CustomerKey AND S.ProductKey = P.ProductKey AND
      P.Discontinued = 'Yes'
GROUP BY C.CustomerName, P.ProductName;
```

Star query

Three basic steps to evaluate the query:

- (1) Evaluation of the join conditions
- (2) Evaluation of the selection conditions over the dimensions
- (3) Aggregation of the tuples that passed the filter

Evaluation of Star Queries with Bitmap Indexes: Example

Product Key	Product Name	...	Discontinued	...
p1	prod1	...	No	...
p2	prod2	...	Yes	...
p3	prod3	...	No	...
p4	prod4	...	Yes	...
p5	prod5	...	No	...
p6	prod6	...	Yes	...

Product table

Yes	No
0	1
1	0
0	1
1	0
0	1
1	0

Bitmap for Discontinued

Customer Key	Customer Name	Address	Postal Code	...
c1	cust1	35 Main St.	7373	...
c2	cust2	Av. Roosevelt 50	1050	...
c3	cust3	Av. Louise 233	1080	...
c4	cust4	Rue Gabrielle	1180	...

Customer table

Product Key	Customer Key	Time Key	Sales Amount
p1	c1	t1	100
p1	c2	t1	100
p2	c2	t2	100
p2	c2	t3	100
p3	c3	t3	100
p4	c3	t4	100
p5	c4	t5	100

Sales fact table

c1	c2	c3	c4
1	0	0	0
0	1	0	0
0	1	0	0
0	1	0	0
0	0	1	0
0	0	1	0
0	0	0	1

Bitmap for
CustomerKey

p1	p2	p3	p4	p5	p6
1	0	0	0	0	0
1	0	0	0	0	0
0	1	0	0	0	0
0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0
0	0	0	0	1	0

Bitmap for
ProductKey

Yes	No
0	1
0	1
1	0
1	0
0	1
1	0
0	1

Bitmap join index for
Discontinued

Evaluation of Star Queries using Bitmap Indexes

Evaluation of star query requires

- a B+ tree over **CustomerKey** and **ProductKey**
- Bitmap indexes on the foreign key columns in **Sales** and on **Discontinued** in **Product**

Example of query evaluation

- (1) Obtain the record numbers of the records that satisfy the condition **Discontinued = 'yes'**
- Answer: Records with **ProductKey** values **p2** , **p4** , and **p6**
- (2) To access the bitmap vectors in **Sales** with these labels perform a join between **Product** and **Sales**
- (3) Vectors labeled **p2** and **p4** match, no fact record for **p6**
- (4) Obtain the values for the **CustomerKey** in these records (**c2** and **c3**)
- (5) Use B+-tree index on **ProductKey** and **CustomerKey** to find the names of products and customers
- (6) Answer: (**cust2, prod2, 200**) and (**cust3, prod4, 100**)

[TOP](#)

Physical Data Warehouse Design

Outline

[Techniques for Physical Data Warehouse Design](#)

[Materialized View](#)

[Indexes for Data Warehouses](#)

[Evaluation of Star Queries](#)

[Data Warehouse Partitioning](#)

Data Warehouse Partitioning

Partitioning (or **fragmentation**) divides a table into smaller data sets (each one called a partition)

Applied to tables and indexes

Vendors provide several different partitioning methods

Vertical partitioning splits the attributes of a table into groups that can be independently stored

- E.g., most often used attributes are stored in one partition, less often used attributes in another one
- More records fit into main memory, reducing their processing time

Horizontal partitioning divides a table into smaller tables with same structure than the full table

- For example, if some queries require the most recent data, partition horizontally according to time

Queries over Partitioned Databases

Partition pruning is the typical way of improving query performance using partitioning

Example: A **Sales** fact table in a warehouse can be partitioned by month

A query requesting orders for a single month only needs to access the partition of such a month

Joins also enhanced by using partitioning:

- When the two tables are partitioned on the join attributes
- When the reference table is partitioned on its primary key
- Large join is broken down into smaller joins

Partitioning Strategies

Three partitioning strategies: Range partitioning, hash partitioning, and list partitioning

Range partitioning maps records to partitions based on ranges of values of the partitioning key

Time dimension is a natural candidate for range partitioning

Example: A table with a **date** column defined as the partitioning key

- **January-2012** partition will contain rows with key values from January 1 to January 31, 2012

Hash partitioning uses a hashing algorithm over the partitioning key to map records to partitions

- Hashing algorithm distributes rows among partitions in a uniform fashion, yielding, ideally, partitions of the same size
- Typically used when partitions are distributed in several devices, and when data are not partitioned based on time

Partitioning Strategies

List partitioning specifies a list of values for the partitioning key

Some vendors (e.g. Oracle) support the notion of **composite partitioning**, combining the basic data distribution methods

Thus, a table can be range partitioned, and each partition can be subdivided using hash partitioning

References

A. VAISMAN, E. ZIMANYI, Data Warehouse Systems: Design and Implementation, Chapter 7 Physical Data Warehouse Design, Springer Verlag, 2014