

ISIT312 Big Data Management

HBase Data Model

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, SIM, Session 4, 2020

HBase Data Model

Outline

Background

Logical view of data

Design fundamentals

Physical implementation

Performance

Background

HBase is open source distributed database based on a data model of Google's **BigTable**

HBase provides a **BigTable** view of data stored in **HDFS**

HBase is also called as **Hadoop DataBase**

HBase still provides a tabular view of data however it is also very different from the traditional **relational data model**

HBase data model is a sparse, distributed, persistent multidimensional sorted map

It is indexed by a **row key**, **column key**, and **timestamp**

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

[Performance](#)

Logical view of data

HBase organizes data into **tables**

HBase **table** consists of **rows**

Each **row** is uniquely identified by a **row key**

Data within a **row** is grouped by a **column family**

Column families have an important impact on the **physical implementation** of **HBase table**

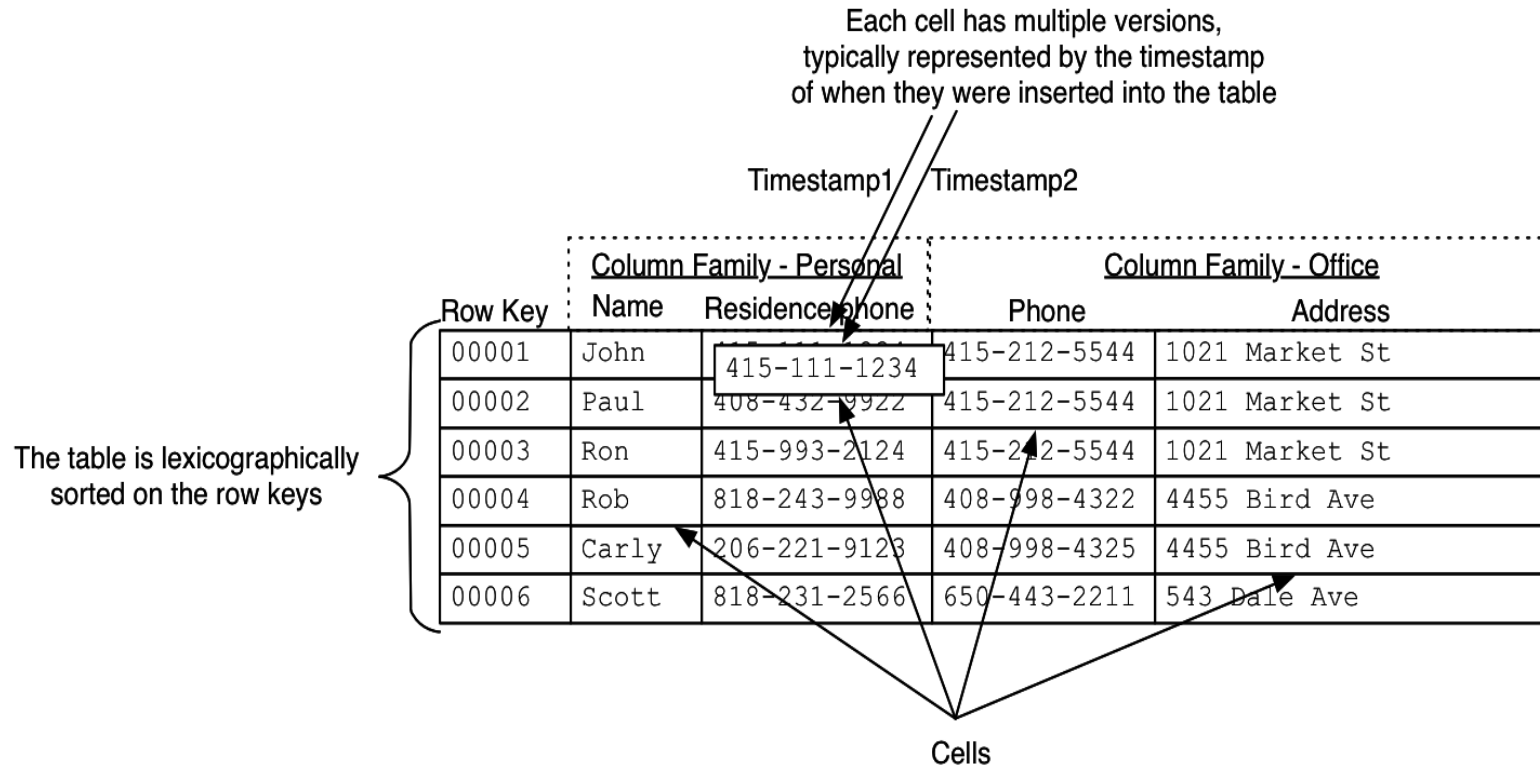
Every **row** has the same **column families** although some of them can be **empty**

Data within a **column family** is addressed via its **column qualifier**, or simply, **column name**

Hence, a combination of **row key**, **column family**, and **column qualifier** uniquely identifies a **cell**

Values in cells do not have a data type and are always treated as **sequences of bytes**

Logical view of data



Values within a cell have multiple versions

Versions are identified by their version number, which by default is a timestamp when the cell was written

Logical view of data

If a **timestamp** is not determined at write time, then the **current timestamp** is used

If a **timestamp** is not determined during a read, the **latest one** is returned

The **maximum allowed number of cell value versions** is determined for each **column family**

The **default number of cell versions** is three

Logical view of data

A view of **HBase table** as a nested structure

```
{ "Row-0001":  
  { "Home":  
    { "Name":  
      { "timestamp-1": "James" }  
      "Phones":  
        { "timestamp-1": "2 42 214339"  
          "timestamp-2": "2 42 213456"  
          "timestamp-3": "+61 2 4567890" }  
        }  
      "Office":  
        { "Phone":  
          { "timestamp-4": "+64 345678" }  
          "Address":  
            { "timestamp-5": "10 Ellenborough Pl" }  
          }  
        }  
      }  
    }  
  }  
}
```

HBase Table

Logical view of data

A view of **HBase table** as a nested structure

```
{ "Row-0002":  
  { "Home":  
    { "Name":  
      { "timestamp-6": "Harry" }  
      "Phones":  
        { "timestamp-7": "2 42 214234" }  
    }  
    "Office":  
      { "Phone":  
        { "timestamp-8": "+64 345678" }  
        "Address":  
          { "timestamp-9": "10 Bong Bong Rd"  
            "timestamp-10": "23 Victoria Rd" }  
        }  
      }  
    }  
  }  
}
```

HBase Table

Logical view of data

A view **HBase table** as a **key->value** store

```
"Row-0001" -> {"Home": {"Name": {"timestamp-1": "James"},
                      "Phones": {"timestamp-2": "2 42 214339",
                                "timestamp-3": "2 42 213456",
                                "timestamp-4": "+61 2 4567890"}},
               "Office": {"Phone": {"timestamp-5": "+64 345678"},
                          "Address": {"timestamp-6": "10 Ellenborough Pl"}},
               }
```

Key->value

```
"Row-0001" "Home" -> {"Name": {"timestamp-1": "James"},
                      "Phones": {"timestamp-2": "2 42 214339",
                                "timestamp-3": "2 42213456",
                                "timestamp-4": "+61 2 4567890"}},
                      }
```

Key->value"

```
"Row-0001" "Home" "Phones" -> {"timestamp-2": "2 42 214339",
                                "timestamp-3": "2 42 213456",
                                "timestamp-4": "+61 2 4567890"}
```

Key->value

Logical view of data

A **key** can be **row key** or a combination of a **row key**, **column family**, **qualifier**, and **timestamp** depending on what supposed to be retrieved

If all the **cells** in a row are of interest then a **key** is a **row key**

If only specific **cells** are of interest, the appropriate **column families** and **qualifiers** are a part of a **key**

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

[Performance](#)

Design Fundamentals

When designing **Hbase table** we have to consider the following questions:

- What should be a **row key** and what should it contain?
- How many **column families** should a **table** have?
- What **columns** (**qualifiers**) should be included in each **column family**?
- What information should go into the **cells**?
- How many **versions** should be stored for each **cell**?

In fact **HBase table** is a four level **hierarchical structure** where a **table** consists of **rows**, **rows** consists of **column families**, **column families** consist of **columns** and **columns** consists of versions

If cells contain the keys then **HBase table** becomes a **network/graph structure**

Design Fundamentals

Important facts to remember:

- **Indexing** is performed only for a **row key**
- **Hbase tables** are stored sorted based on a **row key**
- Everything in **Hbase tables** is stored as **untyped sequence of bytes**
- **Atomicity** is guaranteed only at a row level and there are no multi-row transactions
- **Column families** must be defined at **Hbase table** creation time
- **Column qualifiers** are dynamic and can be defined at write time
- **Column qualifiers** are stored as sequences of bytes such that they can represent data

Design Fundamentals

Implementation of Entity type

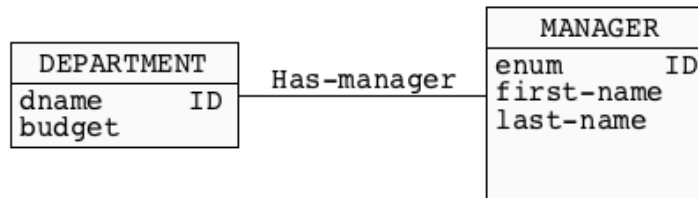
CUSTOMER	
cnumber	ID
first-name	
last-name	
phone	
email	

```
{ "007":  
  { "CUSTOMER":  
    { "first-name": { "timestamp-1": "James" }  
      "last-name": { "timestamp-2": "Bond" }  
      "phone": { "timestamp-1": "007" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of **one-to-one relationship**



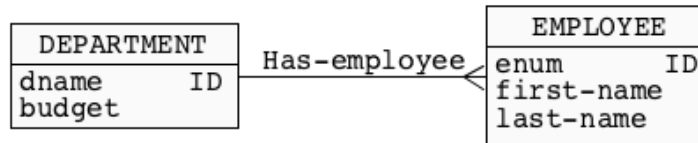
```

{"Sales":
  {"DEPARTMENT":
    {"budget": {"timestamp-1": "1000"}
  }
  {"MANAGER":
    {"enumber": {"timestamp-2": "007"}
    "first-name": {"timestamp-3": "James"}
    "last-name": {"timestamp-4": "Bond"}
  }
}
  
```

HBase Table

Design Fundamentals

Implementation of one-to-many relationship

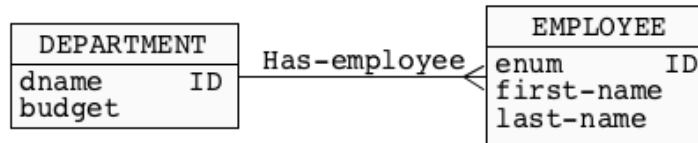


```
{  
  "007":  
    {"EMPLOYEE":  
      {"enum": {"timestamp-1": "007"},  
        "first-name": {"timestamp-2": "James"},  
        "last-name": {"timestamp-3": "Bond"},  
        "department": {"timestamp-4": "Sales"},  
      }  
    }  
}
```

HBase Table

Design Fundamentals

Another implementation of **one-to-many relationship**

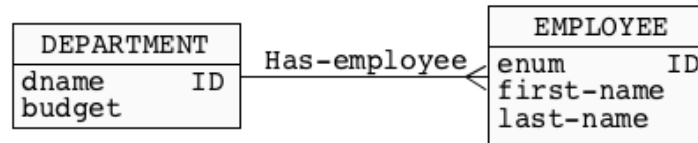


```
{ "Sales":  
  { "DEPARTMENT":  
    { "budget": { "timestamp-1": "1234567" }  
    }  
  }  
  { "EMPLOYEE":  
    { "007": { "timestamp-2": "James Bond" }  
    { "008": { "timestamp-3": "Harry Potter" }  
    { "009": { "timestamp-4": "Robin Hood" }  
    ...  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Yet another implementation of **one-to-many relationship**




```

{"Sales":
  {"DEPARTMENT":
    {"budget": {"timestamp-1": "1000"}}
  }
  {"HAS-EMPLOYEES":
    {"employees": {"timestamp-2": "007 James Bond"}
                  {"timestamp-3": "008 Harry Potter"}
                  {"timestamp-3": "009 Robin Hood"}
                  ...
    }
  }
}
  
```

HBase Table

Design Fundamentals

Implementation of many-to-many relationship




```

{"participation-1":
  {"EMPLOYEE":
    {"enumber": {"timestamp-1": "007"}
    "first-name": {"timestamp-2": "James"}
    "last-name": {"timestamp-3": "Bond"}
    "pnumber": {"timestamp-4": "project-1"}
              {"timestamp-5": "project-2"}
              ...
    }
  }
}
  
```

HBase Table

Design Fundamentals

Another implementation of **many-to-many** relationship




```
{"participation-1":
  {"PROJECT":
    {"pnumber": {"timestamp-1": "project-1"}
    "budget:": {"timestamp-2": "12345.25"}
    "employee": {"timestamp-3": "007"}
                  {"timestamp-4": "008"}
                  {"timestamp-5": "009"}
                  ...
    }
  }
}
```

HBase Table

Design Fundamentals

Another implementation of **many-to-many relationship**



```

{"participation-1":
  {"PARTICIPATION":
    {"pnumber": {"timestamp-1": "project-1"}
    "employee": {"timestamp-2": "employee-007"}
  }
}
  
```

HBase Table

```

{"participation-2":
  {"PARTICIPATION":
    {"pnumber": {"timestamp-1": "project-1"}
    "employee": {"timestamp-2": "employee-008"}
  }
}
  
```

HBase Table

Design Fundamentals


Note, that it is possible to group in one **Hbase table** rows of different types

```
{
  "employee-007": {
    "EMPLOYEE": {
      "enumber": {"timestamp-1": "007"}
      "first-name": {"timestamp-2": "James"}
      "last-name": {"timestamp-3": "Bond"}
    }
  }
  "project-1": {
    "PROJECT": {
      "pnumber": {"timestamp-4": "1"}
      "budget": {"timestamp-5": "12345.25"}
    }
  }
  "participation-2": {
    "PARTICIPATION": {
      "pnumber": {"timestamp-1": "project-1"}
      "employee": {"timestamp-2": "employee-007"}
    }
  }
}
```

HBase Table


Design Fundamentals

Implementation of fact with dimensions



Design Fundamentals

Implementation of fact with dimensions




```
{"1234567":  
  {"MEASURE":  
    {"amount": {"timestamp-1": "1000000"}}  
  },  
  "BUYER":  
    {"phone": {"timestamp-1": "242214339"},  
     "first-name": {"timestamp-1": "James"},  
     "last-name": {"timestamp-1": "Bond"},  
    },  
  "SELLER":  
    {"phone": {"timestamp-1": "242215612"},  
     "first-name": {"timestamp-1": "Harry"},  
     "last-name": {"timestamp-1": "potter"},  
    },  
}
```

HBase Table

Design Fundamentals

Implementation of **graph structure**



```
{ "A":  
  { "R":  
    { "1": { "timestamp-1": "B" }  
      "2": { "timestamp-1": "D" }  
    }  
    "S":  
    { "1": { "timestamp-1": "C" }  
    }  
  }  
}
```


HBase Table

```
{ "B":  
  { "R":  
    { "1": { "timestamp-1": "D" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of **graph structure**



```
{ "C":  
  { "S":  
    { "1": { "timestamp-1": "D" }  
    }  
  }  
}
```

HBase Table

```
{ "D":  
}
```

HBase Table

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

[Performance](#)

Physical implementation

HBase is a database built on top of HDFS

HBase tables can scale up to billions of rows and millions of columns

Because Hbase tables can grow up to terabytes or even petabytes, Hbase tables are split into smaller chunks of data that are distributed across multiple servers

Chunks of data are called as regions and servers that host regions are called as region servers

Region servers are usually colocated with data nodes of HDFS

The splits of Hbase tables are usually horizontal, however, it is also possible to benefit from vertical splits separating column families

Region assignments happen when Hbase table grows in size or when a region server is malfunctioning or when a new region server is added

Physical implementation

Two special HBase tables `-.ROOT-` and `.META.` keep information where the **regions** for the tables are hosted

A table `-.ROOT-` never splits into more than one region, while `.META.` can be distributed over many **regions**

When a client application wants to access a particular row `-.ROOT-` points it to the region of the `.META.` table that contains information about a row location

The entry point for an **HBase** system is provided by another system called **ZooKeeper**

ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

[Performance](#)

Performance

Too many **regions** affect performance

More regions means smaller memory flushes to persistent storage and smaller **HFiles** stored in **HDFS**

It requires **HBase** to process many compaction to keep the number of **HFiles** low

HBase can handle regions from 10 to 40 Gb

Too many regions may occur due to

- Over-splitting with **HBase's split feature**
- Improper presplitting with **HBase presplit feature**

Offline **merging regions** can be used to reduce number of regions

Performance

Too many **column families** affect performance

A magic total number of **column families** number is 3

Column families are built to group data with similar format or similar access pattern

Consequences of too many **column families**

- **Memory**: **HBase** shares write cache among all column families of the same regions, more column families means less transient memory for each one of them
- **Compactions**: the total number of **column families** affects the total number of store files create during flushes and subsequently number of compactions
- **Split**: **HBase** stores **column families** into separate files and directories, when a directory becomes too large it is **split**; **split** affects all **column families**, if some of them are small only few cells are stored in a directory

A solution is a correct schema design followed by **delete column family**, **merge column family**, or **separate column family** in a new **HBase table**

Performance

All writes and reads should be uniformly distributed across the regions

Hotspotting occurs when a given region serviced by a single Region Server receives most or all of the read or write requests

Causes of **hotspotting**:

- **Monotonically incrementing keys**: only last bits or bytes are slowly changing, e.g. timestamp used as a key, all writes to a consecutive area of keys go to the same region
- **Poorly distributed keys**: e.g. keys are always appended with the same prefix going to the same region
- **Small reference (lookup) tables**: All references to lookup tables must be handled by the same region server
- **Application issues**: e.g. writes always performed on the same region
- **Meta region hotspotting**: e.g. creating too many connections

References

Kerzner M., Maniyam S., HBase Design Patterns, Packt Publishing 2014
(Available from UoW Library)

Jiang, Y. HBase Administration Cookbook, Pack Publishing, 2012
(Available from UoW Library)

Dimiduk N., Khurana A., HBase in Action, Manning Publishers, 2013

Spaggiari J-M., O'Dell K., Architecting HBase Applications, O'Reilly, 2016