

ISIT312 Big Data Management

Pig Latin

Dr Guoxin Su and Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Pig Latin

Outline

Preliminary matters

Input/Output in Pig Latin

Operations

Built-in functions

Foreach

Filter

Group

Order by

Distinct

Joins

Limit and sample

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

2/32

Preliminary Matters

Pig Latin is a dataflow language

Each processing step results in a new data set, or relation

Loading, filtering, and transforming

```
A = load 'NYSE_dividends' (exchange, symbol, date, dividends);  
B = filter A by dividends > 0;  
C = foreach B generate UPPER(symbol);
```

Case Sensitivity

- `A = load 'foo';` is not equivalent to `a = load 'foo';`

Comments

Comments

```
A = load 'foo'; --this is a single-line comment  
/*  
    This is a multiline comment.  
*/
```

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

4/32

Input in Pig Latin

load

By default, load looks for your data on **HDFS** in a tab-delimited file using the default **load function PigStorage**

```
divs = load '/data/examples/file';
```

Loading data

A statement above looks for a file called 'file' in a folder
hdfs://[host-name]:[port- name]/data/examples

It is possible to determine a particular **load function**

```
divs = load '/data/examples/file' using HBaseStorage(,);
```

Loading data HBaseStorage load function

It is also possible to specify the schema when loading the file

```
divs = load '/data/examples/file' as (exchange: int, symbol: chararray);
```

Loading with a schema

Output in Pig Latin

store

Pig stores your data on **HDFS** in a tab-delimited file using PigStorage

```
store processed into '/data/examples/processed';
```

Storing data

As in **load**, you can also specify a **store function**, e.g. **HbaseStorage()**

In **processed**, there are usually multiple part files rather than a single file (why ?)

dump

Occasionally you will want to see it on the screen, this is done by the **dump** action

```
dump processed;
```

Displaying data

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

7/32

Arithmetic, Boolean and relational Operations

Pig Latin supports the following basic operations

- arithmetic operations: $+$, $-$, $*$, $/$
- Boolean operations: `and`, `or`, `not`
- relational operations: `==`, `!=`, `<`, `>`, `<=`, `>=`, `matches` (pattern matching)

Relational operations are the main tools Pig Latin provides to operate on your data

They are the horsepower of Pig Latin

They allow you to transform the operations such as `sorting`, `grouping`, `joining`, `projecting`, and `filtering`

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

9/32

Built-In Functions

Mathematical functions

- AVG, COUNT, MAX, MIN, SIZE, SUM, ABS, CEIL, EXP, FLOOR, LOG, RANDOM, ROUND, and others

String functions

- STARTWITH, ENDWITH, LOWER, UPPER, REGEX_EXTRACT, TOKENIZE, and others

Date/time functions

- CurrentTime, DaysBetween, GetDay, GetHour, GetMinute, ToDate, and others

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

11/32

Foreach

foreach takes a set of expressions and applies them to every record in the data pipeline, it is the **projecting** and **transforming** operation

Loading and transforming

```
A = load 'input' as (user:chararray, id:long, address:chararray, phone:chararray,
                    preferences:map[]);
B = foreachA generate user, id;
```

It is convenient to use **expressions** in **foreach** statement

Loading and transforming

```
prices = load 'NYSE_daily' as (exchange, symbol, date, open, high, low, close,
                               volume, adj_close);
gain = foreach prices generate close - open;
gain2 = foreach prices generate $6 - $3;
```

Foreach

Some other expressions

Loading and transforming

```
prices = load 'NYSE_daily' as (exchange, symbol, date, open, high, low,
                                close, volume, adj_close);
beginning = foreach prices generate ..open;
-- produces exchange, symbol, date, open

middle = foreach prices generate open..close;
-- produces open, high, low, close

end = foreach prices generate volume..;
-- produces volume, adj_close

all_in_one = foreach prices generate *;
-- produces a tuple of all fields
```

To extract data from complex types

Extracting data from complex types

```
bball = load 'baseball' as (name:chararray, team:chararray,
                             position:bag{t:(p:chararray)}, bat:map[]);
bball_avg = foreach bball generate bat#'batting_average';
bball_p = foreach bball generate position.p;
```

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

14/32

Filter

filter statement allows you to select which records will be retained in your data pipeline

A filter contains a **predicate**

If a **predicate** evaluates to true for a given record, that record will be passed down the pipeline

Otherwise, it will not

Loading and filtering

```
divs = load 'NYSE_dividends' as (exchange:chararray, symbol:chararray,  
                                date:chararray, dividends:float);  
startswithcm = filter divs by symbol matches 'CM.*';  
-- Only records matching the pattern 'CM.*' are retained
```

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

16/32

Group

group statement collects together records with the same key

It is sharing its syntax with **SQL**, but the grouping operator in **Pig Latin** is fundamentally different than the one in **SQL**

In **SQL GROUP BY** clause creates a group that must feed directly into one or more aggregate functions

In **Pig Latin** there is no direct connection between **group** and aggregate functions

Instead, **group** does exactly what it says, i.e. it collects all records with the same value for the provided key together into a bag

```
daily = load 'NYSE_daily' as (exchange:chararray, stock_id:int);  
grp = group daily by stock;  
describe grp;  
grp: {group: int, daily: {exchange: chararray, stock_id: int}}
```

Loading and grouping

Group

Group supports multiple keys

The resulting records still have two fields and **group** field is a tuple with a field for each key

```
daily = load 'NYSE_daily' as (exchange, stock, date, dividends);
grp = group daily by (exchange, stock);
avg = foreach grp generate group, AVG(daily.dividends);
describe grp;
grp: {group: (exchange: bytearray,stock: bytearray),daily: {exchange: bytearray,
                                                             stock: bytearray,
                                                             date: bytearray,
                                                             dividends: bytearray}}
```

Loading and grouping

It is also possible to use **all** to group together all of the records in a pipeline

```
daily = load 'NYSE_daily' as (exchange, stock);
grp = group daily all;
--grp will have only one row containing all records
```

Loading and grouping

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

19/32

Order By

`order by` statement sorts producing a total order of output data

The syntax of `order by` is similar to `group` statement

Data is sorted based on the types of the indicated fields: **numeric values** are sorted **numerically**, **chararray fields** are sorted **lexically**

Loading and sorting

```
daily = load 'NYSE_daily' as (exchange:chararray, symbol:chararray, date:chararray,  
                                open:float);  
byclose = order daily by close desc, open;  
  
dump byclose;  
-- first close in descending order, then open sorted in ascending order
```

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

21/32

Distinct

`distinct` statement removes duplicate records

```
daily = load 'NYSE_daily' as (exchange:chararray, symbol:chararray);  
uniq = distinct daily;
```

Loading and removing duplicates

`Distinct` statement operates only on entire records and not on individual fields

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

23/32

Joins

`join` operations in **Pig** are analogous to the `join` operations in **SQL**

`join` operations combine records from two bags based on a common field

`join` operations act on two different data sets where one field in each data set is nominated as a **join key**

In a `join`, the **first** (**second**) dataset specified is called as the **left** (**right**) entity or data set

Joins

Types of `join` operation

- `inner join`, often simply called a `join`, returns all elements or records from both datasets where the nominated key is present in both datasets
- `left outer join` returns all records from the `left`, or `first` dataset along with matched records only (by the specified key) from the `right`, or `second`, dataset
- `right outer join` returns all records from the `right`, or `second` dataset along with matched records only (by the specified key) from the `left`, or `first` dataset
- `full outer join` returns all records from both datasets whether there is a key match or not

Inner Join

```
inner = JOIN stores BY storied, salespeople BY storeid;
```

Inner join

stores

```
{  
(100, Hayward),  
(101, Baumholder),  
(102, Alexandria),  
(103, Melbourne)  
}
```

salespeople

```
{  
(1, Henry, 100),  
(2, Karen, 100),  
(3, Paul, 101),  
(4, Jimmy, 102),  
(5, Janice)  
}
```



inner

```
{  
(100, Hayward, 2, Karen, 100),  
(100, Hayward, 1, Henry, 100),  
(101, Baumholder, 3, Paul, 101),  
(102, Alexandria, 4, Jimmy, 102)  
}
```

Left Outer Join

Left outer join

```
leftouter = JOIN stores BY storied LEFT OUTER, salespeople BY storeid;
```

stores

```
{  
(100, Hayward),  
(101, Baumholder),  
(102, Alexandria),  
(103, Melbourne)  
}
```

salespeople

```
{  
(1, Henry, 100),  
(2, Karen, 100),  
(3, Paul, 101),  
(4, Jimmy, 102),  
(5, Janice)  
}
```



leftouter

```
{  
(100, Hayward, 2, Karen, 100),  
(100, Hayward, 1, Henry, 100),  
(101, Baumholder, 3, Paul, 101),  
(102, Alexandria, 4, Jimmy, 102)  
(103, Melbourne, , , )  
}
```

Right Outer Join

Right outer join

```
rightouter = JOIN stores BY storied RIGHT OUTER, salespeople BY storeid;
```

stores

```
{  
(100, Hayward),  
(101, Baumholder),  
(102, Alexandria),  
(103, Melbourne)  
}
```

salespeople

```
{  
(1, Henry, 100),  
(2, Karen, 100),  
(3, Paul, 101),  
(4, Jimmy, 102),  
(5, Janice)  
}
```



rightouter

```
{  
(100, Hayward, 2, Karen, 100),  
(100, Hayward, 1, Henry, 100),  
(101, Baumholder, 3, Paul, 101),  
(102, Alexandria, 4, Jimmy, 102)  
(, , 5, Janice , )  
}
```

Full Outer Join

Full outer join

```
fullouter = JOIN stores BY storied FULL OUTER, salespeople BY storeid;
```

stores

```
{  
(100, Hayward),  
(101, Baumholder),  
(102, Alexandria),  
(103, Melbourne)  
}
```

salespeople

```
{  
(1, Henry, 100),  
(2, Karen, 100),  
(3, Paul, 101),  
(4, Jimmy, 102),  
(5, Janice)  
}
```



fullouter

```
{  
(100, Hayward, 2, Karen, 100),  
(100, Hayward, 1, Henry, 100),  
(101, Baumholder, 3, Paul, 101),  
(102, Alexandria, 4, Jimmy, 102)  
(103, Melbourne, , , )  
( , , 5, Janice, )  
}
```

Pig Latin

Outline

[Preliminary matters](#)

[Input/Output in Pig Latin](#)

[Operations](#)

[Built-in functions](#)

[Foreach](#)

[Filter](#)

[Group](#)

[Order by](#)

[Distinct](#)

[Joins](#)

[Limit and sample](#)

[TOP](#)

Created by Janusz R. Getta, ISIT312 Big Data Management, Session 2, 2021

30/32

Limit and Sample

Sometimes we want to see only a limited number of results

limit statement allows for the restriction of an output to a given number of the first records

The following returns at most 10 lines, i.e., the first 10 records

```
divs = load 'NYSE_dividends';  
first10 = limit divs 10;
```

Loading and displaying limited number of records

Sometimes we want to see only a limited number of randomly selected results

sample statement allows for the restriction of an output to a given percentage of the total number of records

In the following example, 0.1 indicates 10%

```
divs = load 'NYSE_dividends';  
some = sample divs 0.1;
```

Loading and displaying sample number of records

References

Gates A., Programming Pig, O'Reilly Media, Inc., 2011, (Available in **READINGS** folder)

Vaddeman B., Beginning Apache Pig: Big Data Processing Made Easy, Apress 2016 (Available in **READINGS** folder)